Support Vector Machines

Quiz, 5 questions

Support Vector Machines

1
point

1.

Suppose you have trained an SVM classifier with a Gaussian kernel, and it learned the following decision boundary on the training set:

You suspect that the SVM is underfitting your dataset.

Should you try increasing or decreasing C? Increasing or decreasing σ^2?

○ It would be reasonable to try **decreasing** C. It would also be reasonable to try **increasing** σ^2.

○ It would be reasonable to try **increasing** C. It would also be reasonable to try **decreasing** σ^2.

○ It would be reasonable to try **increasing** C. It

Support Vector Machines

1
point

2.

The formula for the Gaussian kernel is given by

$$\text{similarity}(x, l^{(1)}) = \exp\left(-\frac{||x - l^{(1)}||^2}{2\sigma^2}\right).$$

The figure below shows a plot of $f_1 = \text{similarity}(x, l^{(1)})$ when

$\sigma^2 = 1.$

Which of the following is a plot of f_1 when $\sigma^2 = 0.25$?

Support Vector Machines

1
point

3.

The SVM solves

$$\min_\theta \ C \sum_{i\,=\,1}^m y^{(i)}\text{cost}_1(\theta^T x^{(i)}) + (1 - y^{(i)})\text{cost}_0(\theta^T x^{(i)}) + \sum_{j\,=\,1}^n \theta_j^2$$

where the functions $\text{cost}_0(z)$ and $\text{cost}_1(z)$ look like this:

The first term in the objective is:

$$C \sum_{i\,=\,1}^m y^{(i)}\text{cost}_1(\theta^T x^{(i)}) + (1 - y^{(i)})\text{cost}_0(\theta^T x^{(i)}).$$

This first term will be zero if two of the following four conditions hold true. Which are the two conditions that would guarantee that this term equals zero?

☐   For every example with $y^{(i)} = 0$, we have that

$\theta^T x^{(i)} \leq 0$.

☐   For every example with $y^{(i)} = 1$, we have that

$\theta^T x^{(i)} \geq 1$.

☐   For every example with $y^{(i)} = 0$, we have that

Support Vector Machines

4.

Suppose you have a dataset with n = 10 features and m = 5000 examples.

After training your logistic regression classifier with gradient descent, you find that it has underfit the training set and does not achieve the desired performance on the training or cross validation sets.

Which of the following might be promising steps to take? Check all that apply.

- [ ] Increase the regularization parameter λ.

- [ ] Use an SVM with a Gaussian Kernel.

- [ ] Use an SVM with a linear kernel, without introducing new features.

- [ ] Create / add new polynomial features.

Support Vector Machines

| 1 |
| point |

5.

Which of the following statements are true? Check all that apply.

- [ ] If the data are linearly separable, an SVM using a linear kernel will

  return the same parameters θ regardless of the chosen value of

  C (i.e., the resulting value of θ does not depend

  on C).

- [ ] It is important to perform feature normalization before using the Gaussian kernel.

- [ ] Suppose you are using SVMs to do multi-class classification and

  would like to use the one-vs-all approach. If you

  have K different

  classes, you will train K - 1 different SVMs.

- [ ] The maximum value of the Gaussian kernel (i.e.,

  $sim(x, l^{(1)})$) is 1.

- [ ] I, **Saprem Mukesh Shah**, understand that submitting work that isn't my own may result in permanent failure of this course or deactivation of my Coursera account.
  Learn more about Coursera's Honor Code

Submit Quiz

Support Vector Machines

Quiz, 5 questions